

Introduction to R

Mark van der Loo

UFPEL, October 2019

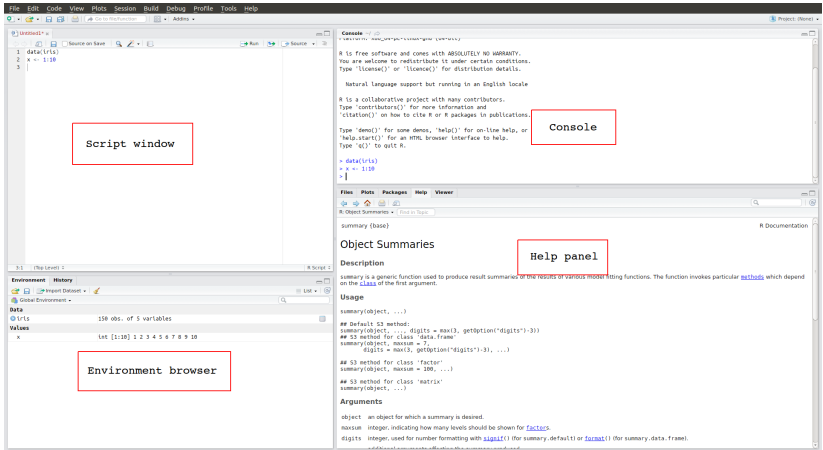
Contents, today's goal

- ▶ R, RStudio
- ▶ Working with command-line and scripts
- ▶ Basic data types (vector, data.frame)
- ▶ Some plots
- ▶ Basic data processing

R and RStudio

R and the R community





Console

- ▶ Connects to the 'R interpreter'
- ▶ You can type commands there or copy them from the script window
- ▶ Resultats are printed to the console again.

```
1 + 1
```

```
[1] 2
```

Script window


- ▶ Here you can open and edit several types of text files, e.g.
 - ▶ .R (R scripts)
 - ▶ .Rmd to create reports that include your results
 - ▶ C/C++ for programming with C or C++
- ▶ Use CTRL-ENTER to send the currently selected command to the R interpreter.
- ▶ The script window is the single most important place in RStudio! **WRITE ALL YOUR CODE IN SCRIPTS.**

Environment browser

- ▶ Gives an interactive overview of all data loaded into R
 - ▶ data sets, results of modeling; anything really.
- ▶ You can get the same overview by typing `ls()` in the command-line

Help panel

- Help pages for each R function



- Open a help page for a function: ?<function> or search: ??<search term>.

Note

The help pages are pretty dense and technical. They are aimed to be technical documentation, but don't be intimidated! There is lots of help online.

Getting help

- ▶ Q-and-A site `stackoverflow.com`
 - ▶ Easily found via Google.
 - ▶ n00b-friendly
- ▶ R-help mailinglist `r-project.org/mail.html`
 - ▶ You may get answers from the R-core developers.
 - ▶ DO READ THE POSTING GUIDE

Tip of the day

Error message? Cut-and-paste it in Google.

Literature

- ▶ Working with R:
 - ▶ R in a Nutshell (J. Addler) *O'Reilly*
 - ▶ **R for data science** (H. Wickham and G. Grolemund) *O'Reilly*
- ▶ Programming, package development:
 - ▶ The Art of R Programming (N. Matloff) *No Starch Press*
 - ▶ Testing R code (R. Cotton) *O'Reilly*
 - ▶ R Packages (H. Wickham) *O'Reilly*
 - ▶ Advanced R (H. Wickham) *CRC Press*
- ▶ Applications:
 - ▶ *Use R!* series: www.springer.com/series/6991
 - ▶ *The R Series* crcpress.com/go/the-r-series
 - ▶ ...
- ▶ See also r-project.org/doc/bib/R-books.html

Basic data types and the R command-line

Some tips

Repeat commands

Use arrow keys \uparrow , \downarrow to cycle through previous commands

Keyboard shortcuts (in script window)

CTRL+ENTER Execute current command

CTRL-SHIFT-S Execute current script

Auto-complete


Use tab to complete names of objects, columns in `data.frames` and file names (between quotes).

Vectors

The basic unit in R is a *vector*: a sequence of values of the same type (like a column of data in SAS or SPSS –but not Excel!).

Example

```
# Sample 100 numbers from the normal distribution  
# Store under the name 'x'  
x <- rnorm(100)  
# plot a histogram of x  
hist(x, breaks=20)
```



Example (cont'd) statistical summaries.

```
summary(x) # overview
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-2.57753	-0.54664	0.06918	0.12463	0.78946	2.85021

```
sd(x) # standard deviation
```

```
[1] 1.070897
```

```
head(x,3) # first three values
```

```
[1] -0.30295450 -0.16087374 0.09127928
```


Example (cont'd) metadata

```
length(x)
```

```
[1] 100
```

```
class(x)
```

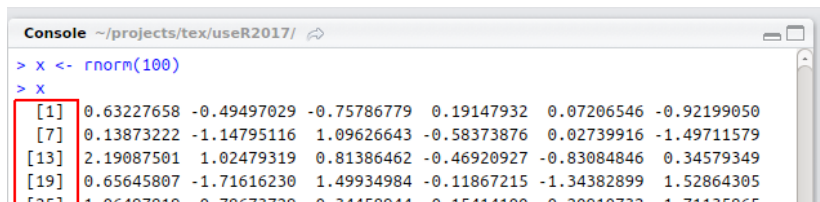
```
[1] "numeric"
```

```
y <- c(joe=1, bill=7, averett=3)  
names(y)
```

```
[1] "joe"      "bill"     "averett"
```

Some observations

- ▶ You can create and name vectors under (almost) any name. Use `<-` to store something under a given name.
- ▶ You do calculations with *functions*, like `sd`, `min`, `mean`
- ▶ When a vector is printed, the first column in the terminal shows the position.



The screenshot shows a terminal window titled "Console ~/projects/tex/useR2017/". The user has entered the R command `x <- rnorm(100)` and then `x` to print the vector. The output displays the first few elements of the vector, with the first column indicating their positions. A red box highlights the first four rows of the output, showing the indices [1], [7], [13], and [19].

```
> x <- rnorm(100)
> x
[1] 0.63227658 -0.49497029 -0.75786779 0.19147932 0.07206546 -0.92199050
[7] 0.13873222 -1.14795116 1.09626643 -0.58373876 0.02739916 -1.49711579
[13] 2.19087501 1.02479319 0.81386462 -0.46920927 -0.83084846 0.34579349
[19] 0.65645807 -1.71616230 1.49934984 -0.11867215 -1.34382899 1.52864305
[25] 1.06107010 0.70673720 0.34450044 0.15111100 0.20010732 1.71135065
```

Creating vectors

<code>c(...)</code>	Assign value by value (<code>x <- c(1,6,2)</code>)
<code>seq(from, to, [by])</code>	Create a sequence (<code>x <- seq(1,10,2)</code>)
<code>seq_len(length.out)</code>	Create a sequence <code>1,2,...,length.out</code>
<code>:</code> (dubbele punt)	<code>a:b</code> gives <code>a,a+1,...,b</code>
<code>rnorm(n, [mean], [sd])</code>	Sample from normal distribution
<code>runif(n, [min], [max])</code>	Sample from uniform distribution

Opmerkingen

- ▶ Argumenten in square brackets are optional.
- ▶ `seq()` also works for time/data sequences

Summarizing vectors

<code>mean,median</code>	mean, median
<code>sum</code>	Sum
<code>min,max</code>	Minimum, maximum
<code>sd</code>	Standaard deviation
<code>fivenum</code>	Tukey's five-number statistics
<code>summary</code>	Sammary (works for all types)
<code>hist</code>	Histogram
<code>boxplot</code>	Boxplot
<code>length</code>	Nr of elements in a vector
<code>class</code>	Type of data
<code>names</code>	Labels

Remeber that

R is case sensitive

```
x <- 10  
X <- 11  
ls()
```

```
[1] "x" "X" "y"
```

Variabelen can be overwritten

```
x <- 10  
x <- "fiets"  
x
```

```
[1] "fiets"
```

Computing with vectors

Addition etc works element-by-element.

```
x <- c(1,3,2,6)
```

```
y <- c(2,5,7,3)
```

```
x + y # add
```

```
[1] 3 8 9 9
```

```
x * y # multiply
```

```
[1] 2 15 14 18
```

```
x ^ y # x to the power of y
```

```
[1] 1 243 128 216
```

Computing with vectors (cont'd): Recycling

For vectors of unequal length, the shorter is repeated

```
x
```

```
[1] 1 3 2 6
```

```
2 * x  # here is '2' a vector of length 1
```

```
[1] 2 6 4 12
```

```
x + 3
```

```
[1] 4 6 5 9
```

Transformations

All the usual math functions are available

```
x <- c(0,1,4,9, 12)
sqrt(x) # square root of x
```

```
[1] 0.000000 1.000000 2.000000 3.000000 3.464102
```

Examples

exp, log, log10

sqrt

sin, cos, tan, sinh, cosh, tanh

Data types

numeric	Numbers (integer or real)
integer	Integers
logical	Boolean (TRUE,FALSE)
character	Text
factor	Categorical (nominal) data
POSIXct	Date/time

Opmerkingen

- ▶ R converts automatically from integer to numeric
- ▶ There are a few more types (complex, raw) not shown here

Missing values

- ▶ Missing values are represented with NA.
- ▶ Almost any calculation involving NA will result in NA

```
x <- c(1,4,2,NA,6)
c( mean1 = mean(x), mean2 = mean(x, na.rm=TRUE) )
```

```
mean1 mean2
NA    3.25
```

- ▶ Skip NA with na.rm=TRUE

RStudio project | data import | data frames

Contents

- ▶ Create an RStudio project
- ▶ Scripts
- ▶ Reading csv files
- ▶ Introducing dplyr

Reading text files with readr

Reading

<code>read.csv</code>	Comma for columns, dot for decimals
<code>read.csv2</code>	Semicolon for columns, comma for decimals
<code>read.table</code>	Any 'rectangular' text data.

Wegschrijven

<code>write.csv</code>	Kommascheiding, punt is decimaalteken
<code>write.csv2</code>	Punktkommascheiding, komma is decimaalteken
<code>write.table</code>	Alle rechthoekige bestanden in tekstformaat.

```
dat <- read.csv("myfile.csv")  
write.csv2(dat, "yourfile.csv", row.names=FALSE)
```

File names in R

- ▶ Always in quotes.
- ▶ It can also be a url.
- ▶ Always use forward slash as directory separator:

```
dat <- read.csv("C:/users/joe/documents/foo.csv")
```

Tip of the day

Always work in an RStudio project. It makes it much easier to locate files.

Data frames

Een data.frame is a bunch of vectors of the same length.

this dataset is built into R for examples.

```
head(InsectSprays,3)
```

	count	spray
1	10	A
2	7	A
3	20	A

Summarizing data frames

```
summary(InsectSprays)
```


	count	spray
Min.	: 0.00	A:12
1st Qu.:	3.00	B:12
Median	: 7.00	C:12
Mean	: 9.50	D:12
3rd Qu.:	14.25	E:12
Max.	:26.00	F:12

Some handy functions

Funcitie	description
<code>summary</code>	Statistical summary
<code>str</code>	Technical summary
<code>colMeans, rowMeans</code>	mean per column, row
<code>colSums, rowSums</code>	sum per column, row
<code>names</code>	column names
<code>ncol nrow</code>	nr of columns, rows
<code>dim</code>	vector with <code>nrow</code> , <code>ncol</code>


Plotting (1)

```
plot(weight ~ height, data=women)
```



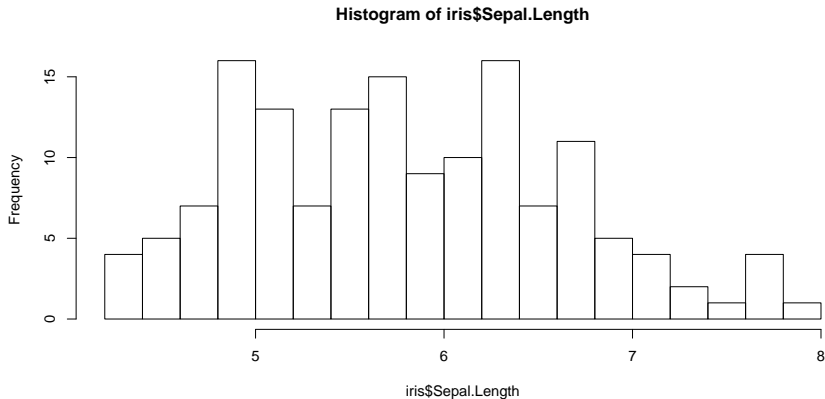
Plotting (2)

```
plot(count ~ spray, data=InsectSprays)
```



Plotting (3)

```
# met '$' selecteer je een kolom  
hist(iris$Sepal.Length, breaks=20)
```



Introduction to data manipulation with dplyr

```
library(dplyr)
```

Verbs for common operations

filter	Rijen selecteren
select	Kolommen selecteren
rename	Kolommen hernoemen
distinct	Unieke rijen selecteren
arrange	Sorteren
transmute	Nieuwe kolommen berekenen
mutate	Nieuwe kolommen toevoegen

dplyr::filter

Select rows.

```
filter(.data, ...)
```

Here, .data is a data.frame (or tibble) and ... are conditions.

```
filter(iris, Sepal.Length > 7)
```

```
filter(iris, Sepal.Length > 7, Species=="virginica")
```

```
filter(iris, Sepal.Length > mean(Sepal.Length))
```

Comparison operators

Expression	TRUE when
<code>x == y</code>	x equals y
<code>x <= y</code>	x does not exceed y
<code>x < y</code>	x strictly smaller than y
<code>x > y</code>	x strictly larger than y
<code>x >= y</code>	x larger than or equal to y
<code>x != y</code>	x unequal to y
<code>x %in% y</code>	x appears in y

Example: %in%

```
x <- c("noot", "boom", "roos", "vis", "aap")  
y <- c("aap", "noot", "mies")  
x %in% y
```

```
[1] TRUE FALSE FALSE FALSE TRUE
```


Logical operators

Operator	Betekenis
&	AND
	OR (en/of)
!	NOT
all(x)	are all entries in x TRUE?
any(x)	is at least entry in x TRUE?

dplyr::select

Select columns

```
select(.data, ...)
```

Use ... to select columns:

```
select(iris, Sepal.Width, Petal.Width)
```

Or give the selected columns new names:

```
select(iris, bladlengte=Petal.Length  
      , soort=Species)
```

dplyr::rename

Rename columns

```
rename(.data, ...)
```

Specify as <new name> = <old name>.

```
rename(iris, species = Species)  
rename(iris, leaf_size = Sepal.Width, species=Species)
```

dplyr::distinct

Keep only unique rows

```
distinct(.data, ..., .keep_all=FALSE)
```

With ... you specify what columns determine whether a record is unique. In case of duplicates, the first record is kept. The keep_all option determines whether to keep all columns or just the ones specified in

```
distinct(iris, Species, keep_all=TRUE)
```

dplyr::arrange

Sorteer de rijen.

```
arrange(.data, ...)
```

Use ... to specify sorting variables. Each next variable is a tie-breaker for the previous ones. Use desc to sort descending instead of increasing.

```
arrange(iris, Sepal.Length, Petal.Width)  
arrange(iris, Sepal.Length, desc(Petal.Width))
```

dplyr::mutate

Add columns

```
mutate(.data, ...)
```

Use ... to specify a sequence of expressions that define the new columns.

```
mutate(women  
  , lengthM    = height * 2.54/100  
  , weightKg   = weight/2.046  
  , bmi        = weightKg/(lengthM^2))
```

Expressions are always in the form <new name> = <expression>.

dplyr::transmute

Compute new columns

```
transmute(.data, ...)
```

Same as mutate, except only the new columns are returned.

```
transmute(women, ratio=height/weight)
```